

SVD Factorization for Tall-and-fat Matrices on Map/Reduce Architectures

Burak Bayramlı

October 15, 2013

Abstract

We demonstrate an implementation for an approximate rank- k SVD factorization, combining well-known randomized projection techniques with previously implemented map/reduce solutions in order to compute steps of the random projection based SVD procedure, such as QR and SVD. We structure the problem in a way that it reduces to Cholesky and SVD factorizations on $k \times k$ matrices computed on a single machine, greatly easing the computability of the problem.

1 Introduction

[1] presents many excellent techniques for utilizing map/reduce architectures to compute QR and SVD for the so-called tall-and-skinny matrices. The idea is based on the fact that QR factorization can be turned into an $A^T A$ computation problem computed in parallel, en masse using map/reduce, and through this to a Cholesky decomposition performed on a single machine. Since

$$A^T A = (QR)^T (QR) = R^T Q^T QR = R^T R$$

and because Cholesky factorization of an $n \times n$ symmetric positive definite matrix is

$$A = LL^T$$

where L is an $n \times n$ lower triangular matrix, and R is upper triangular, we can conclude if we factorize A into L and L^T , this implies $LL^T = RR^T$, we have a method of calculating R of QR using Cholesky factorization on $A^T A$. The key observation here is $A^T A$ computation results in an $n \times n$ matrix and

if A is “skinny” then n is relatively small (in the thousands), and Cholesky decomposition can be executed on a small $n \times n$ matrix on a single computer. Q is computed simply as $Q = AR^{-1}$. This again is relatively cheap because R is $n \times n$, the inverse is computed locally, matrix multiplication with A can be performed through map/reduce.

SVD is an additional step. SVD decomposition is

$$A = U\Sigma V^T$$

If we expand it with $A = QR$

$$QR = U\Sigma V^T$$

$$R = Q^T U \Sigma V^T$$

Let's call $\tilde{U} = Q^T U$

$$R = \tilde{U} \Sigma V^T$$

This means if we run a local SVD on R (we just calculated above with Cholesky) which is an $n \times n$ matrix, we will have calculated \tilde{U} , the real Σ , and real V^T .

Now we have a map/reduce way of calculating QR and SVD on $m \times n$ matrices where n is small.

1.1 Approximate rank-k SVD

Switching gears, we look at another method for calculating SVD. The motivation is computing SVD if n is large, creating a “fat” matrix which might have columns in the billions would require reducing the dimensionality of the problem. According to [2], one way to achieve is through random projection. First we draw an $n \times k$ Gaussian random matrix Ω . Then we calculate

$$Y = A\Omega$$

We perform QR decomposition on Y

$$Y = QR$$

Then form $k \times n$ matrix

$$B = Q^T A$$

Then we can calculate SVD on this small matrix

$$B = \hat{U}\Sigma V^T$$

Then form the matrix

$$U = Q\hat{U}$$

The main idea is based on

$$A = QQ^T A$$

if replace Q which comes from random projection Y ,

$$A \approx \tilde{Q}\tilde{Q}^T A$$

Q and R of the projection are close to that of A . In the multiplication above R is called B where $B = \tilde{Q}^T A$, and,

$$A \approx \tilde{Q}B$$

then, as in [1], we can take SVD of B and apply the same transition rules to obtain an approximate U of A .

This approximation works because of the fact that projecting points to a random subspace preserves distances between points, or in detail, projecting the n -point subset onto a random subspace of $O(\log n/\epsilon^2)$ dimensions only changes the interpoint distances by $(1 \pm \epsilon)$ with positive probability [3]. It is also said that Y is a good representation of the span of A .

1.2 Combining Both Methods

What if n is also very large? In this case local Cholesky or SVD computations would take a long time as well. Our idea was using approximate k -rank SVD where $k \ll n$, before map/reduce based QR and SVD methods presented in section 1, to reduce dimension before using map/reduce methods presented in Section 1, this way, we are again able to work with small matrices locally, $k \times k$ this time on which Cholesky, SVD can be performed. Below we outline each map/reduce job.

Algorithm 1: Random Projection Job

```
input : A
output: Y
function MAP(key, value)
    Tokenize value and pick out id value pairs
    result  $\leftarrow$  zeros(1,k)
    for each  $j^{th}$  token  $\in$  value do
        Initialize seed with j
        r  $\leftarrow$  generate k random numbers
        result  $\leftarrow$  result + r  $\cdot$  token[j]
    end
    emit key, result
function REDUCE(key, value)
    noop
```

Each value of A will arrive to the algorithm as a key and value pair. Key is line number or other identifier per row of A . Value is a collection of id value pairs where id is column id this time, and value is the value for that column. Sparsity is handled through this format, if an id for a column does not appear in a row of A , it is assumed to be zero. The resulting Y matrix has dimensions $m \times k$.

Algorithm 2: $A^T A$ Cholesky Job

```
input : Y
output: R
function MAP(key k, val a)
    for i, row in enumerate( $a^T a$ ) do
        | emit i, row
    end
function REDUCE(key, value)
    | emit (k, sum( $\langle v_j^k \rangle$ ))
function FINAL LOCAL REDUCE (key, value)
    | result  $\leftarrow$  Cholesky( $A_{sum}$ )
    | emit (result)
```

The FINAL_LOCAL_REDUCE step is a function provided in most map/reduce frameworks, it is a central point that collects the output of all reducers, naturally a single machine which makes it ideal to execute the final Cholesky call on a $k \times k$ matrix. The output is R .

Algorithm 3: Q Job

```
input :  $Y, R$ 
output:  $Q$ 
function  $INIT()$ 
|  $R_{inv} = R^{-1}$ 
function  $MAP(key, value)$ 
| for  $row$  in  $Y$  do
| | emit ( $key, row \cdot R_{inv}$ )
| end
```

There is no reducer in the Q Job, it is a very simple job, it merely computes multiplication between row of Y and a local matrix R . Matrix R is very small, $k \times k$, hence it can be kept locally in every node. The $INIT$ function is used to store the inverse of R locally, once the mapper is initialized, it will always use the same R^{-1} for every multiplication.

Algorithm 4: $A^T Q$ Job

```
input :  $A, Q$ 
output:  $B^T$ 
function  $REDUCE(key, value)$ 
| for  $row$  in  $value$  do
| | if  $row$  is from  $A$  then
| | |  $left = row$ 
| | end
| | if  $row$  is from  $Q$  then
| | |  $right = row$ 
| | end
| end
| for  $nonzero\ j^{th}\ cell\ in\ left$  do
| | emit  $j, left[j] \cdot right$ 
| end
function  $REDUCESUM(key, value)$ 
|  $result \leftarrow zeros(1, k)$ 
| for  $row$  in  $value$  do
| |  $result \leftarrow result + row$ 
| end
| emit  $key, result$ 
```

The job above takes A and Q matrices at the same time. Both of these matrices are based on the same key (line number, of preexisting id) and we need to join them first. If a mapper is a pass-through mapper, in other words if it does not exist, it is assumed to simply re-emit the key and value,

which will, indirectly force matching rows with the same id from A and Q to be sent to the first reducer. Then for each unique id, the first reducer gets exactly two rows. This is an indirect way of performing a join in map/reduce environment.

Then in the reducer we deduce if the row is a Q row or an A row. We need this information because we will iterate cells of A one by one, which is assumed to be sparse. In this iteration for each j^{th} non-zero cell of A , we multiply the cell's value with the row from Q and emit the multiplication result with key j .

References

- [1] Gleich, Benson, Demmel, *Direct QR factorizations for tall-and-skinny matrices in MapReduce architectures*
- [2] N. Halko, *Randomized methods for computing low-rank approximations of matrices*
- [3] S. Dangupta, A. Gupta *An Elementary Proof of a Theorem of Johnson and Lindenstrauss*